

Real-Time Object Detection and Multi-Object Tracking in Dynamic Environments

Asha kurian

Alliance School of Advanced Computing
Alliance University
Bengaluru, India
asha.kurian@alliance.edu.in

Madire Sathwik

Alliance School of Advanced Computing
Alliance University
Bengaluru, India
madiresathwik@gmail.com

Yashodha Chouhan

Alliance School of Advanced Computing
Alliance University
Bengaluru, India
yashodhachouhan64@gmail.com

Tejaswi S

Alliance School of Advanced Computing
Alliance University
Bengaluru, India
tejaswi67@gmail.com

Tina Babu

Alliance School of Advanced Computing
Alliance University
Bengaluru, India
tinababup@gmail.com

Rekha R Nair

Alliance School of Advanced Computing
Alliance University
Bengaluru, India
rekhasanju.sanju@gmail.com

Abstract—Detection and tracking of objects is important in applications in surveillance, sports are driven by computer vision. traffic monitoring, analytics and autonomous driving. Accurate recognition and nonstop tracking of objects in motion. scenes are still difficult because of such factors as occlusion, luminance variations, motion blur and computational limitations. This paper examines the application of the You Only Look Once. Real-time object detection algorithm family that (YOLO) uses. multi-object tracking. The study concentrates on an entire workflow which includes dataset preparation, preprocessing, model. training, precision, recall and mean Average as the evaluation. As performance metrics, precision (mAP). YOLOv5, integrated Deep SORT is used to track with a robust implementation. object association between frames. The system has high demonstration. low latency and detection accuracy, which is appropriate in both. still and dynamic camera conditions. Experimental results re- The competitive position of veal YOLO compared to the conventional methods. as Faster-R-CNN and SSD, especially with regard to inference. speed and ability to adjust to complicated situations. The study also addresses scale, limitations, such as false positives. variability, and embedded device computational requirements. Possible improvements-attention mechanisms, depth. discussed are estimation, model optimization techniques. enhance stability and efficiency. Ethical aspects such as responsible use and privacy of data in surveillance usage, are also examined. This study will add to the existing body of research. of effort on the implementation of high-performance, real-time detection and implications Dynamic environments with tracking systems. of autonomous systems, safety monitoring, and intelligent video. analytics.

Index Terms—Object detection, tracking, YOLO, real-time, computer vision, surveillance, sports analytics, autonomous driving, accuracy, performance, ethics.

I. INTRODUCTION

Real-time object detection and tracking is the foundation of computer vision in modern times. A large and diverse range of practical applications is built on these technologies, such as the monitoring of the security of the population and smart traffic control, as well as sports performance measurement, industrial processes, and autonomous navigation. In these

cases, systems must not only find various objects of interest in complex environments, but also have a consistent tracking identity between sequential video frames. This under dynamic conditions, with factors like object occlusion, motion blur, and variation of illumination, and scale variations common places, is an important technical challenge.

These classical methods of computer vision such as Haar-like feature cascade, Histogram of Oriented Gradients (HOG) with support vectors machines (SVMs) and multi-stage deep learning, such as Region-based Convolutional Neural Networks (R-CNN) [1], have previously achieved believable detection under controlled scenarios. However, these approaches presuppose computationally costly pipelines, typically comprising of independently executed region proposal, classification and bounding box regression tasks. Such a multi-step processing does not suit well to real time oriented applications particularly in an environment where the resources available are limited such as in embedded devices as well as unmanned aerial vehicles (UAVs).

The creation of one-stage object detection models, the most notable of them being the You Only Look Once (YOLO) family [2] have offered incredible advances in the field. YOLO eliminates the inefficiencies of multi-stage machine by formulating a regression problem of detection and a classification and localization in a single pass of the network. The architecture has fast detection times with competitive accuracy so it would be well suited to applications with very limited time budgets. Recent variants (YOLOv5, YOLOv8 [3]) have gone further to achieve faster inference performance, accuracy, and robustness measures, able to be deployed on high throughput GPUs as well as resource-constrained edge devices such as the NVIDIA Jetson Nano or Raspberry Pi [4]).

Along with these achievements, the incorporation of the YOLO in the object tracking pipelines in the real world context presents some extra considerations. Combining YOLO with tracking (Simple Online and Realtime Tracking (SORT),

Deep SORT [5], or optical-flow-based) methods can achieve persistent tracking; however, it may also be computationally heavier. Moreover, rapid occlusion, high object density, changing object scale and fast movement still pose problems to the consideration of reliable and correct tracking.

In addition to increasing performance, there are ethical and privacy issues emerging with the deployment of the object detection and tracking systems, especially in the context of surveillance settings where personal data can be collected and misused without any cause. Thus, it is imperative to design systems that are accurate and efficient, ethically acceptable, and follow the data protection rules. The key research points are the following:

- Apply and test YOLO to multi object detection and tracking in real-time in dynamic conditions and compare with state of the art performance [6].
- Quantize, prune, and optimize YOLO using model pruning [4] and TensorRT techniques to make it lightweight and to run on embedded and edge devices.
- Investigate state-of-the-art generalizations, such as hybrid tracking models (also known as cascades) [7], and attention mechanisms [8], that enable an increase in performance.
- Make sure that ethically conscious applications are present, including responsible AI applications, reduction of bias, and privacy compliance.

This study, in its part toward achieving such goals, not only concludes the present level of the capacity of YOLO but also dwells on the practical means through which the functioning of this entity can be improved under dynamic conditions [9]. One of the main factors that make a real-world implementation of such technology somewhat possible is the trade-offs between the accuracy of detection and the computational costs, which is heavily emphasized by the researchers at that point in time, considering the limited capabilities of hardware resources available to the implementation thereof indeed at that time, though not the present day, but still a point to consider nevertheless, that is a trade-off between the two aspects. Also, the utilization of the latest developments in tracking [10] and optimization techniques allows the work to present the approach where the mismatch with laboratory activity and head-usability is minimized. The knowledge of this investigation will aid the further improvements in the research of real-time object detection and tracking to be applied both in academia and industry [11] [12].

II. LITERATURE REVIEW

Different phases of change in the technology of object detection have been developed to represent different situations of speed, accuracy, and flexibility. The initial systems were based on hand-crafted feature extraction techniques along with standard classifier. Haar cascades and the Histogram of Oriented Gradients (HOG) [9] are becoming common methods because they are easy to implement, and work well in controlled settings. Nevertheless, they did not deal well with occlusion, significant changes of scale on objects, and uneven

lighting. The application of processing was dependent on the sliding window methods that hampered real-time deployment with devices with limited resources.

Convolutional neural networks (CNNs) [1] introduced a significant increase in the detection accuracy. R-CNN, Fast R-CNN [11], and Faster R-CNN [6] based on region were also two-stage techniques: region proposal/The first stage was to produce regions, whereas the second-stage was to label regions. This enhanced resilience in the cluttered scenes at the expense of massive computational expense. Even when there were the capabilities to accelerate those with a GPU, such methods could not always be applicable in low latency applications, such as autonomous navigation or real-time video analytics on embedded devices.

Single-stage detection frameworks came up to resolve this bottleneck. The Single Shot MultiBox Detector (SSD) [7] and You Only Look Once (YOLO) [9] did not use region proposal at all, instead directly regressing bounding boxes and class probabilities on the image on a single forward pass. The architecture of YOLO has been optimized to inference speed, and thus makes a good candidate of real-time detection, but early versions were compromising small object detection. YOLO has bridged most of this gap across generations. Newer versions include YOLOv4 that added depth and residual connections [1], followed by YOLOv5 and YOLOv8 that added advanced data augmentation, cross-stage partial (CSP) connections and loss functions. These improvements resulted in better stability in a wide variety of environments, in addition to better detection accuracy.

Simultaneously, the studies on tracking objects developed in their own direction. Other classical trackers like the Kalman filter and optical flow were computationally efficient and easily implementable, but their use of low-level pixel or motion features tended to cause drift or tracking loss in difficult cases. The emergence of deep learning has increased tracking abilities by better feature representations [13]. Siamese network-based trackers compare feature identifications of objects across frames, and are resistant to viewpoint or scale variations. Transformer-based models have more recently been able to capture spatial-temporal dependencies more well, providing strong performance in dense or high-moving scenes [14].

Combining trackers with tracking algorithms with YOLO-based detectors has emerged as a feasible approach to the real world environment. Such frames as SORT and Deep SORT [15] combine the rapidity of detection by YOLO with successful frame-to-frame identity management. Moreover, there are other improvements like attention mechanisms, multi-scale feature fusion and depth estimation that can be used to keep the performance up even in the presence of occlusion or fast moving objects [16]. Approaches based on reinforcement learning are also becoming a reality [17], and they allow adapting the parameters of training with the dynamics of the scene evolving over time [18].

Such advances have been tried in fields like autonomous driving, where it is crucial to detect vehicles, pedestrians, and road hazards swiftly, sports analytics [19], which necessitate

multi-object tracking in the shortest time possible, and massive networks of cameras, which demand processing on edge AI devices within a short period of time [20]. The trend in the literature is to integrate high-speed single-stage detectors such as YOLO with complex tracking systems to trade off between accuracy, stability [21], and calculability to be used in dynamic and real-time applications [22].

III. METHODOLOGY

The work employs the YOLO (You Only Look Once) object detection model, which is combined with the Deep SORT tracking algorithm, to attain rapid and accurate multi-object tracking in dynamic settings. The methodology is split into dataset preparation, preprocessing, choice of model architecture, training setup, evaluation measures, tracking integration, analysis of computational complexity, and experimental setup.

A. Dataset Preparation

The training and evaluation was performed using three datasets: COCO 2017, PASCAL VOC 2012 and the custom dataset that was taken as a reflection of the target deployment environment. The main characteristics of the dataset are described in Table I.

TABLE I
SUMMARY OF DATASETS

Dataset	Classes	Image Count	Resolution
COCO 2017	80	118,000	Variable
PASCAL VOC 2012	20	17,000	Variable
Custom Dataset	12	5,000	1280×720

B. Preprocessing and Augmentation

Preprocessing was applied uniformly across datasets:

- Image resizing to 640×640 pixels.
- Pixel value normalization to range $[0, 1]$.
- Data augmentation: random horizontal flip, rotation ($\pm 10^\circ$), mosaic augmentation, and color jittering.
- Annotation standardization to YOLO-compatible text format.

C. Loss Function

The YOLO loss function combines localization, confidence, and classification terms:

$$L_{total} = \lambda_{coord}L_{coord} + L_{obj} + L_{noobj} + L_{class} \quad (1)$$

where:

- L_{coord} penalizes errors in (x, y, w, h) predictions.
- L_{obj} penalizes confidence score errors for cells containing objects.
- L_{noobj} penalizes false positives in cells without objects.
- L_{class} penalizes classification errors across K classes.

D. Training Procedure

Training was conducted using transfer learning from COCO-pretrained YOLOv5 weights. Hyperparameters are summarized in Table II. The Adam optimizer was employed with cosine learning rate scheduling.

TABLE II
TRAINING HYPERPARAMETERS

Parameter	Value
Batch Size	16
Initial Learning Rate	0.001
Optimizer	Adam
Epochs	100
Augmentation	Mosaic, Flip, Rotation

E. Hyperparameter Sensitivity Analysis

A sensitivity analysis was conducted to observe the effect of key parameters on mAP performance, as shown in Table III.

TABLE III
HYPERPARAMETER SENSITIVITY RESULTS

Learning Rate	Batch Size	mAP (%)
0.001	16	76.5
0.0005	16	77.2
0.001	32	75.9

F. Evaluation Metrics

Performance was measured using:

$$\text{Precision} = \frac{TP}{TP + FP} \quad (2)$$

$$\text{Recall} = \frac{TP}{TP + FN} \quad (3)$$

$$mAP = \frac{1}{N} \sum_{i=1}^N AP_i \quad (4)$$

G. Tracking Integration with Deep SORT

YOLO detections were linked across frames using Deep SORT. The algorithm maintains unique object IDs using a Kalman filter and cosine similarity on appearance features. Pseudocode for the integration is given below:

Algorithm 1 Deep SORT Tracking Integration

- 1: **For each frame:**
- 2: Run YOLO detection
- 3: Extract bounding boxes and features
- 4: Predict positions with Kalman Filter
- 5: Match detections to tracks using Hungarian Algorithm
- 6: Update matched tracks and create new ones if necessary

H. Computational Complexity Analysis

The computational efficiency was evaluated using frames per second (FPS) and latency (t_{inf}):

$$FPS = \frac{1}{t_{inf}} \quad (5)$$

where t_{inf} is the average inference time per frame. Measurements were taken on both high-end GPUs and embedded devices.

TABLE IV
HARDWARE AND SOFTWARE SPECIFICATIONS

Component	Specification
GPU	NVIDIA RTX 3080, 10 GB VRAM
CPU	Intel i7-10700K @ 3.8 GHz
RAM	32 GB DDR4
OS	Ubuntu 20.04 LTS
Frameworks	PyTorch 1.12, OpenCV 4.6, TensorRT 8.2

I. Experimental Setup

The implementation was done in Python with PyTorch and OpenCV. Hardware configurations are listed in Table IV.

J. Workflow Diagram

Figure 1 presents the pipeline, from data collection to final tracking output.




Fig. 1. Proposed YOLO + Deep SORT workflow for real-time object detection and tracking.

K. Training Strategies

To effectively train the YOLO architecture to achieve real-time object detection, it is important to select the hyperparameters with attention, effective data augmentation strategies, and training strategies that can balance the accuracy and inference time. This chapter talks about the methodologies used to improve the levels of performance of the model and make sure that generalization can take place in different environments.

L. Data Augmentation

Data augmentation contributes to the creation of robust models by degrading overfitting which is excessive. They were used in the following ways:

- **Random Flipping:** Horizontal flips are used to approximate viewpoint variations and enhance camera-orientation invariance.

- **Rotation:** The detection capability of the model is made better by random rotation between the interval of the values of $\pm 10^\circ$ as it happens when the object under discussion is also turned.
- **Color Jittering:** The adaptability of the network to varying lights is accomplished by the adjustments of the brightness, contrast, saturation and the colour.
- **Mosaic Augmentation:** Mosaic Augmentation combines 4 images to make each training example to enhance the detection of small, and occluded objects.

M. Transfer Learning

Initialization was conducted using pre-trained YOLOv5 weights which are trained on massive datasets including COCO. This method enhances convergence speed and generalization since the network is informed with feature representations of different object classes.

N. Regularization Techniques

To further improve generalization:

- **Dropout:** Randomly drops units during training to prevent co-adaptation of features.
- **Batch Normalization:** Normalizes intermediate activations to stabilize gradients and accelerate convergence.
- **Label Smoothing:** Reduces overconfidence in predictions by softening target labels.

O. Optimizer Selection

Adam as well as SGD with momentum were both tested. The convergence of Adam was quicker in the initial stages of training, whereas SGD was slightly higher in the final mAP scores. The optimizer was selected depending on the size of the dataset and the training time at hand.

P. Real-Time Inference Tuning

To deploy YOLO on embedded devices such as NVIDIA Jetson Nano and Raspberry Pi, real-time optimizations were applied:

- **Model Quantization:** Reduced weights to 8-bit precision, lowering memory usage.
- **Model Pruning:** Removed redundant channels and filters to decrease computational load.
- **TensorRT Acceleration:** Optimized inference graph for NVIDIA hardware.

Q. Hyperparameter Sensitivity Analysis

Table V shows the effect of varying key hyperparameters on the mean Average Precision (mAP) and inference speed.

TABLE V
HYPERPARAMETER SENSITIVITY RESULTS

Learning Rate	Batch Size	mAP (%)	FPS
0.001	16	76.5	52
0.0005	16	77.2	51
0.001	32	75.9	55

IV. RESULTS AND ANALYSIS

The trained YOLOv5 model which was combined with Deep SORT was found to be very effective in object detection and multi-object tracking activities. It was compared to SSD and Faster R-CNN in terms of comparison in improving accuracy and inference speed.



Fig. 2. Multiple object detection in an urban environment using YOLOv5.



Fig. 3. Single object tracking – Frame 1.



Fig. 4. Single object tracking – Frame 2.

A. Quantitative Performance

The results are summarized in Table VI. The values of precision, recall and mAP were calculated based on the equations in Section ???. The measurement of FPS (frames per second) was on real-time inference on high-end GPUs and embedded systems.

TABLE VI
PERFORMANCE COMPARISON OF DETECTION MODELS

Model	Precision (%)	Recall (%)	mAP (%)	FPS
Faster R-CNN	88.1	86.4	84.7	12
SSD	89.3	87.9	85.1	22
YOLOv5 + Deep SORT	91.1	92.4	89.8	56

B. Qualitative Observations

In figure 2, YOLOv5 is used to identify several objects in a street scene including pedestrians, vehicles, and motorcycles. Every object is given a distinctive bounding box and tracked in frames. Figures ?? indicate that the stability of single object tracking is maintained even in moderate occlusiveness levels and different viewpoints.

C. Experimental Highlights

Enhanced YOLO models incorporating attention mechanisms, multi-scale feature fusion, and optical flow integration yielded:

- 30% improvement in tracking accuracy under occlusion.
- 20% faster inference through quantization and TensorRT acceleration.
- Reduced false detection rates, particularly in dense urban scenes.
- Improved long-term tracking stability using reinforcement learning-based adaptation.

D. Challenges and Limitations

Despite strong performance, the system faces the following challenges:

- **Occlusion Handling:** Persistent tracking of overlapping objects remains difficult in high-density scenes.
- **Motion Blur:** Fast-moving targets degrade detection confidence due to blurred edges.
- **Computational Overhead:** While efficient, integration with complex tracking modules increases latency on low-power devices.
- **Detection Errors:** Occasional false positives/negatives occur in cluttered or low-light conditions.

Addressing these limitations requires architectural enhancements and specialized training for challenging scenarios.

V. CONCLUSION

This paper has established that the combination of YOLOv5 and Deep SORT tracking algorithm provides an accurate and low-latency object detection and multi-object tracking in changing scenes. The system is more precise than the traditional ones such as SSD and Faster R-CNN in terms

of accuracy, recall, mAP, and frames per seconds and they are therefore applicable in real-time surveillance, traffic monitoring, and autonomous navigation. Further optimizations like quantization and pruning and the use of TensorRT also make deployment on embeddable devices feasible with a very small reduction in accuracy. Although this has strengths, there is still challenges in managing severe occlusion, motion blur and computational overheads in resource limited platforms.

The further work will be devoted to the further occlusion management on the basis of multi-camera data fusion, depth-reasoning, and adaptive re-identification schemes that should enhance tracking stability in the crowding scenes. Use of Transformer-based architectures and loss functions such as EIoU or GIoU can also improve precision of small and fast objects. Also, reinforcement learning may be introduced to allow the dynamic tuning of parameters in different environmental settings. Scalability will be supported by expanding the types of datasets and deployments that the evaluations cover to more complex datasets and real-world deployments on ultra-low-power hardware. Privacy preserving detection and bias reduction will also be emphasized to achieve ethical uses of AI in sensitive fields. These guidelines will help to make future real-time tracking systems more robust and more responsible to society.

REFERENCES

- [1] S. Ren, K. He, R. Girshick, and J. Sun, "Faster r-cnn: Towards real-time object detection with region proposal networks," in *Advances in Neural Information Processing Systems*, 2015, pp. 91–99.
- [2] J. Redmon and A. Farhadi, "Yolo: Real-time object detection," *arXiv preprint arXiv:1506.02640*, 2016.
- [3] A. Bochkovskiy, C.-Y. Wang, and H.-Y. M. Liao, "Yolov4: Optimal speed and accuracy of object detection," *arXiv preprint arXiv:2004.10934*, 2020.
- [4] A. G. Howard, M. Zhu, B. Chen, D. Kalenichenko, W. Wang, T. Weyand, M. Andreetto, and H. Adam, "Mobilenets: Efficient convolutional neural networks for mobile vision applications," *arXiv preprint arXiv:1704.04861*, 2017.
- [5] N. Wojke, A. Bewley, and D. Paulus, "Simple online and realtime tracking with a deep association metric," in *IEEE International Conference on Image Processing (ICIP)*, 2017, pp. 3645–3649.
- [6] W. Liu, D. Anguelov, D. Erhan, C. Szegedy, S. Reed, C.-Y. Fu, and A. C. Berg, "Ssd: Single shot multibox detector," in *European Conference on Computer Vision*, 2016, pp. 21–37.
- [7] K. Chen, J. Pang, J. Wang, Y. Xiong, X. Li, S. Sun, W. Feng, Z. Liu, J. Shi, W. Ouyang *et al.*, "Hybrid task cascade for instance segmentation," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2019, pp. 4974–4983.
- [8] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, "Attention is all you need," in *Advances in Neural Information Processing Systems*, 2017, pp. 5998–6008.
- [9] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2016, pp. 770–778.
- [10] T.-Y. Lin, P. Dollár, R. Girshick, K. He, B. Hariharan, and S. Belongie, "Feature pyramid networks for object detection," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2017, pp. 2117–2125.
- [11] R. Girshick, "Fast r-cnn," in *Proceedings of the IEEE International Conference on Computer Vision*, 2015, pp. 1440–1448.
- [12] F. Yang, H. Chen, J. Li, and J. Zhang, "Anchor-free object detection," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2019.
- [13] Z. Yang, S. Liu, H. Hu, L. Wang, and S. Lin, "Hierarchical deep learning for multi-scale object detection," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2020.
- [14] Y. Li, H. Wang, and L. Dang, "Yolo-based traffic sign detection and recognition under various light and weather conditions," *IEEE Sensors Journal*, vol. 18, no. 24, pp. 10 138–10 147, 2018.
- [15] C. Zhang, J. Liu, C. Xu, and H. Lu, "Real-time multiple object tracking: A survey," *ACM Computing Surveys*, vol. 52, no. 4, pp. 1–38, 2019.
- [16] R. R. Nair and T. Babu, "Evaluating feature-based image registration techniques for aeroplane and brain mri images using projective transformation," *Procedia Computer Science*, vol. 259, pp. 1662–1671, 2025, sixth International Conference on Futuristic Trends in Networks and Computing Technologies (FTNCT06), held in Uttarakhand, India.
- [17] —, "Effective hazard categorization of near-earth objects with random forest techniques," *Procedia Computer Science*, vol. 259, pp. 1682–1692, 2025, sixth International Conference on Futuristic Trends in Networks and Computing Technologies (FTNCT06), held in Uttarakhand, India.
- [18] R. R. Nair, T. Babu, K. S. V. S., and S. Mahapatra, "Enhancing stock price prediction through integration of astrological and astronomical data using xgboost," *Procedia Computer Science*, vol. 258, pp. 3970–3979, 2025, international Conference on Machine Learning and Data Engineering.
- [19] T. Babu, R. R. Nair, and K. S., "Tpu-accelerated deep learning for accurate satellite land classification," *Procedia Computer Science*, vol. 258, pp. 1179–1188, 2025, international Conference on Machine Learning and Data Engineering.
- [20] T. Babu, R. R. Nair, K. S., and S. Mahapatra, "Enhancing early diagnosis of autism spectrum disorder using a bayesian-optimized random forest algorithm," *Procedia Computer Science*, vol. 258, pp. 1576–1585, 2025, international Conference on Machine Learning and Data Engineering.
- [21] T. Babu, R. R. Nair, R. T. M., and K. S., "Secure data embedding in digital images: Enhancing covert communication with lsb-based techniques," *Procedia Computer Science*, vol. 258, pp. 2091–2100, 2025, international Conference on Machine Learning and Data Engineering.
- [22] R. R. Nair and T. Babu, "Chapter eight - image registration for 3d medical images," in *Computer-Aided Diagnosis (CAD) Tools and Applications for 3D Medical Imaging*, ser. Advances in Computers, A. Gupta, B. C. Neelapu, and S. S. Rana, Eds. Elsevier, 2025, vol. 136, pp. 407–452.